

Tutorial 9 Transaction (PostgreSQL version)

Designed by [ZHU Yueming](#)

Descriptions words of basic concepts of ACID in this tutorial are borrowed from lecture notes of CHENG Jingde.

Experimental Objective

- Learn the properties of ACID of transaction.
- Learn three different isolation levels of transaction.

Preparing Your Lab

- Create two consoles in datagrip named Session1 and Session2.
- Execute following queries in either Session1 or Session2.

```
create table students(  
    id serial primary key ,  
    name varchar(30) not null,  
    grade int,  
    unique (name)  
);
```

ACID

1. Atomicity

- The “all or nothing” property.
- A transaction is an indivisible unit that is either performed in its entirety or is not performed at all.
- It is the responsibility of the recovery subsystem of the DBMS to ensure atomicity.

```
begin;  
insert into students (name, grade) values ('name1',90);  
insert into students(name,grade) values('name2',85);  
commit;
```

We only can execute all statements or null statements in one transaction.

2. Consistency

- A transaction must transform the database from one consistent state to another consistent state.
- It is the responsibility of both the DBMS and the application developers to ensure consistency.

```
begin;
update students set grade=grade-4 where name='name1';
update students set grade=grade+4 where name='name2';
commit;
```

If we want to give 4 grades to the second student from the first student, we could not only minus 4 grade from the first student without adding 4 grade to the second in one transaction.

3. Durability

- The effects of a successfully completed (committed) transaction are permanently recorded in the database and must not be lost because of a subsequent failure.
- It is the responsibility of the recovery subsystem to ensure durability.

```
begin ;
truncate students; commit;
rollback ;
select * from students;
```

In a transaction, if we remove all students, and then we commit changes. After that, we could not find any record in students table, even though we execute rollback.

4. Isolation

- Transactions execute independently of one another.
- The partial effects of incomplete transactions should not be visible to other transactions.
- It is the responsibility of the concurrency control subsystem to ensure isolation.

More specific examples of isolation properties could be shown in following description.

Basic commands of transaction

- Change auto commit off (by command line)

```
\set AUTOCOMMIT off
```

- show auto commit (by command line)

```
\set
```

- show current transaction isolation:

```
show transaction_isolation;
```

- set transaction isolation level in current transaction:

```
set transaction isolation level [isolation level];
```

- set default transaction isolation level:

```
set default_transaction_isolation = '[isolation level]';
```

Isolation Levels

As show in the table below, the isolation level of different reading results is only suitable for postgresQL. In fact, there are only three distinct isolation levels in postgresQL, which correspond to the levels read committed, repeatable read and serializable read. Read uncommitted do the same work as read committed in postgresQL, which is different from Mysql or other DBMS.

ISOLATION LEVEL	Dirty Read	Nonrepeatable Read	Phantom Read
read uncommitted	N	Y	Y
read committed	N	Y	Y
repeatable read	N	N	N
serializable	N	N	N

1. Read-uncommitted

One transaction can read the result of query executed by another uncommitted transaction, but it is not supported by postgresQL.

2. Read committed

Nonrepeatable read occurs in step4 and step6.

Step	Session 1	Session 2
1	<code>set default_transaction_isolation = 'read committed';</code>	<code>set default_transaction_isolation = 'read committed';</code>
2	<code>begin;</code>	<code>begin;</code>
3	<code>insert into students(name, grade) values ('name1',90);</code>	
4		<code>select * from students;</code>
5	<code>commit;</code>	
6		<code>select * from students;</code>
7		<code>commit;</code>

Phantom Read occurs in step 6.

Step	Session 1	Session 2
1	<code>set default_transaction_isolation = 'read committed';</code>	<code>set default_transaction_isolation = 'read committed';</code>
2	<code>begin;</code>	<code>begin;</code>
3		<code>insert into students(name, grade) values ('name2',80);</code>
4		<code>commit;</code>
5	<code>update students set grade=100 where grade<100;</code>	
6	<code>select * from students;</code>	
7	<code>commit;</code>	
8	<code>select * from students;</code>	

3. Repeatable read

Nonrepeatable Read and Phantom Read don't occur in step 5 and step 7 respectively.

Only when committed in session 2, the changes in other session can be readed .

Step	Session 1	Session 2
1	<code>set default_transaction_isolation = 'repeatable read';</code>	<code>set default_transaction_isolation = 'repeatable read';</code>
2	<code>begin;</code>	<code>begin;</code>
3	<code>insert into students(name, grade) values ('name2',55);</code>	
4	<code>commit;</code>	
5		<code>select * from students;</code>
6		<code>update students set grade=100 where grade<100;</code>
7		<code>select * from students;</code>
8		<code>commit;</code>
9		<code>select * from students;</code>

4. serializable

Read/write should be serializable executed in different transactions.

An error will occurred in step 8:

[40001] ERROR: could not serialize access due to read/write dependencies among transactions
 Detail: Reason code: Canceled on identification as a pivot, during write. Hint: The transaction might succeed if retried.

Step	Session 1	Session 2
1	<code>set default_transaction_isolation = 'serializable';</code>	<code>set default_transaction_isolation = 'serializable';</code>
2	<code>begin;</code>	<code>begin;</code>
3	<code>insert into students(name, grade) values ('name3',80);</code>	
4	<code>select * from students;</code>	
5		<code>insert into students(name, grade) values ('name4',80);</code>
6	<code>commit;</code>	
7		<code>select * from students;</code>
8		<code>commit;</code>